


# What is Snowflake?

# What is Snowflake?



# Data Platform



## Data Warehouse

Structured & relational data

ANSI Standard SQL

ACID compliant transactions

Data stored in databases, schemas & tables



## Data Lake

Scalable storage and compute

Schema does not need to be defined upfront

Native processing of semi-structured data formats



## Data Engineering

COPY INTO & Snowpipe

Separate compute clusters

Tasks and Streams

All data encrypted at rest and in transit.



## Data Science

Remove data management roadblocks with centralised storage

Partner eco-system includes data science tooling:

- Amazon SageMaker
- DataRobot
- Dataiku



## Data Sharing

Secure Data Sharing

Data Marketplace

Data Exchange

BI with the Snowflake partner ecosystem tools



## Data Applications

Connectors and Drivers

UDFs and Stored Procedures

External UDFs

Preview features such as Snowpark

# Cloud Native



Snowflake's software is purpose built for the Cloud.

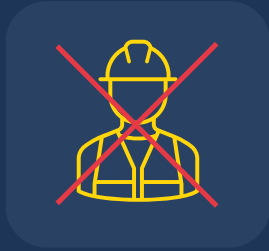


All Snowflake infrastructure runs on the Cloud in either AWS, GCP or Azure.



Snowflake makes use of Cloud's elasticity, scalability, high availability, cost-efficiency & durability.

# Software as a service (SaaS)



No management of  
hardware



Transparent updates  
and patches



Subscription payment  
model



Ease of access

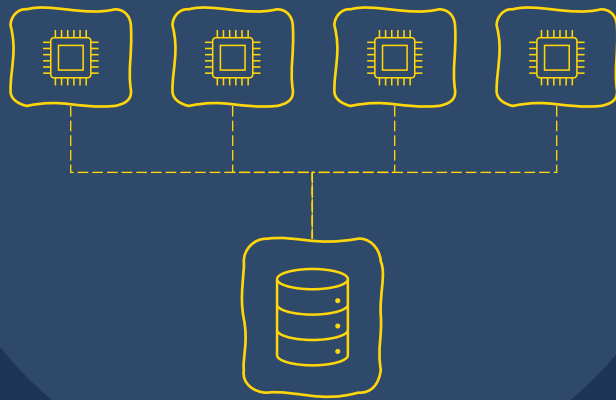


Automatic optimisation

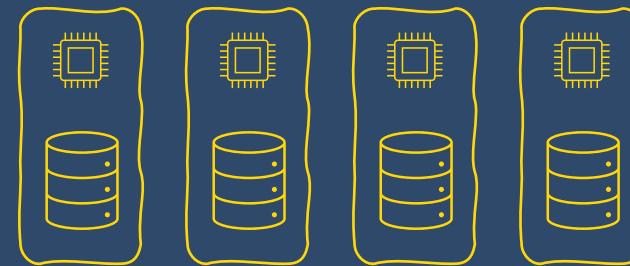
# Multi-cluster Shared Data Architecture

# Distributed Architectures

Shared-Disk




Shared Nothing



# Shared Disk Architecture

Shared-Disk



## Advantages

- Relatively simple to manage
- Single source of truth

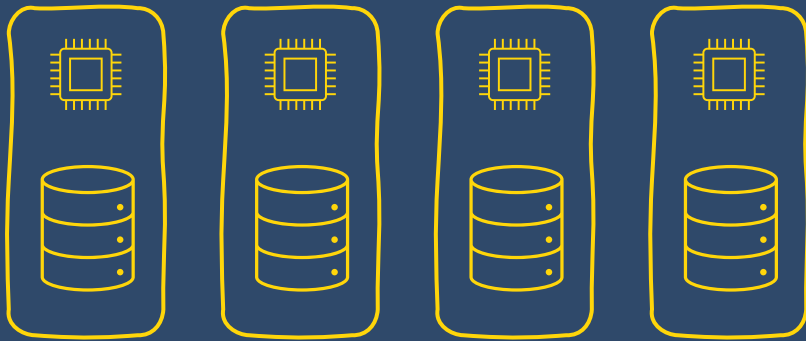
## Disadvantages

- Single point of failure
- Bandwidth and network latency
- Limited scalability



# Shared Nothing Architecture

## Shared Nothing



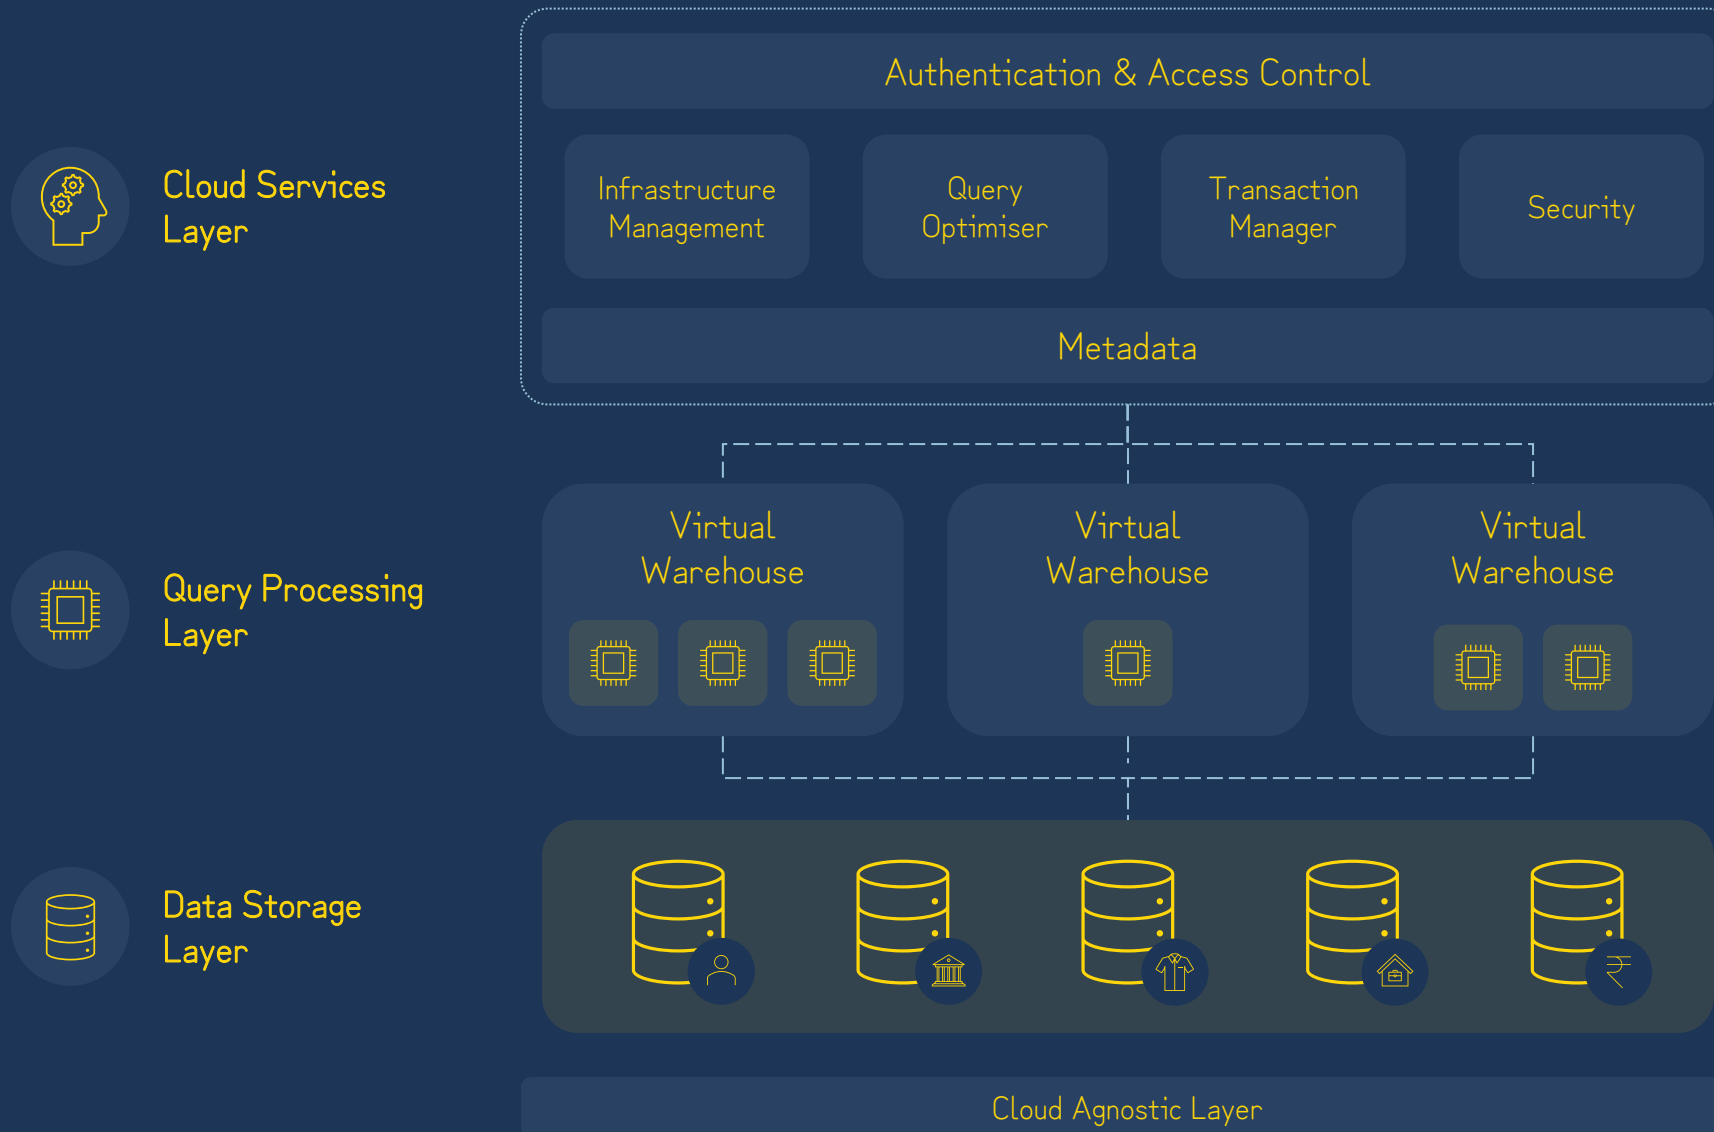
### Advantages

- Co-locating compute and storage avoids networking latency issues
- Generally cheaper to build & maintain
- Improved scaling over shared-disk architecture

### Disadvantages

- Scaling still limited
- Storage and compute tightly coupled
- Tendency to overprovision

# Multi-cluster Shared Data Architecture



- ✓ Decouple storage, compute and management services.
- ✓ Three infinitely scalable layers.
- ✓ Workload isolation with virtual warehouses.

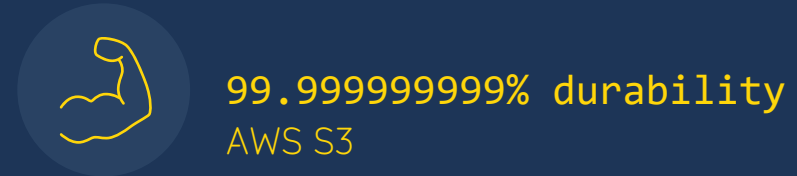
# Storage Layer

# Storage Layer

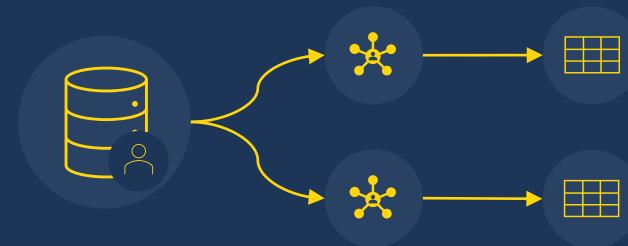
Persistent and infinitely scalable cloud storage residing in cloud providers blob storage service, such as AWS S3.



Snowflake users by proxy get the availability & durability guarantees of the cloud providers blob storage.



Data loaded into Snowflake is organized by databases, schemas and accessible primarily as tables.



Both structured and semi-structured data files can be loaded and stored in Snowflake.



# Storage Layer

When data files are loaded or rows inserted into a table, Snowflake reorganizes the data into its proprietary compressed, columnar table file format.



The data that is loaded or inserted is also partitioned into what Snowflake call micro-partitions.



Storage is billed by how much is stored based on a flat rate per TB calculated monthly.



**\$42.00(TB/mo)**

AWS Europe London

Data is not directly accessible in the underlying blob storage, only via SQL commands.



**SELECT \* FROM <table>;**

**TOM BAILEY COURSES**

[tombaileycourses.com](https://tombaileycourses.com)

# Query Processing Layer


# Query Processing Layer

The query processing layer consists of “Virtual Warehouses” that execute the processing tasks required to return results for most SQL statements.

A **Virtual Warehouse** is a named abstraction for a cluster of a cloud-based compute instances that Snowflake manage.

```
CREATE WAREHOUSE MY_WH WAREHOUSE_SIZE=LARGE;
```

Underlying nodes of a Virtual Warehouse cooperate in a similar way to a shared-nothing compute clusters making use of local caching.



# Query Processing Layer

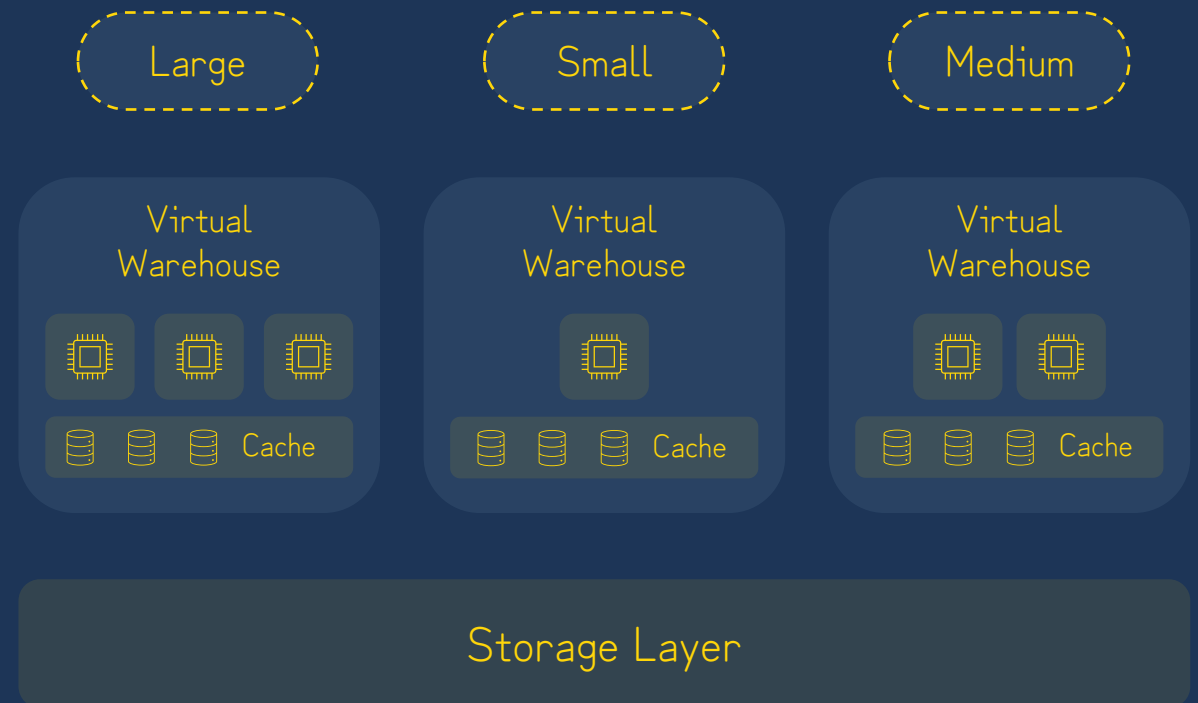
Virtual warehouses can be created or removed instantly.

Virtual warehouses can be paused or resumed.

Virtually unlimited number of virtual warehouses can be created each with it's own configuration.

Virtual warehouses come in multiple “t-shirt” sizes indicating their relative compute power.

All running virtual warehouses have consistent access to the same data in the storage layer.





# Services Layer

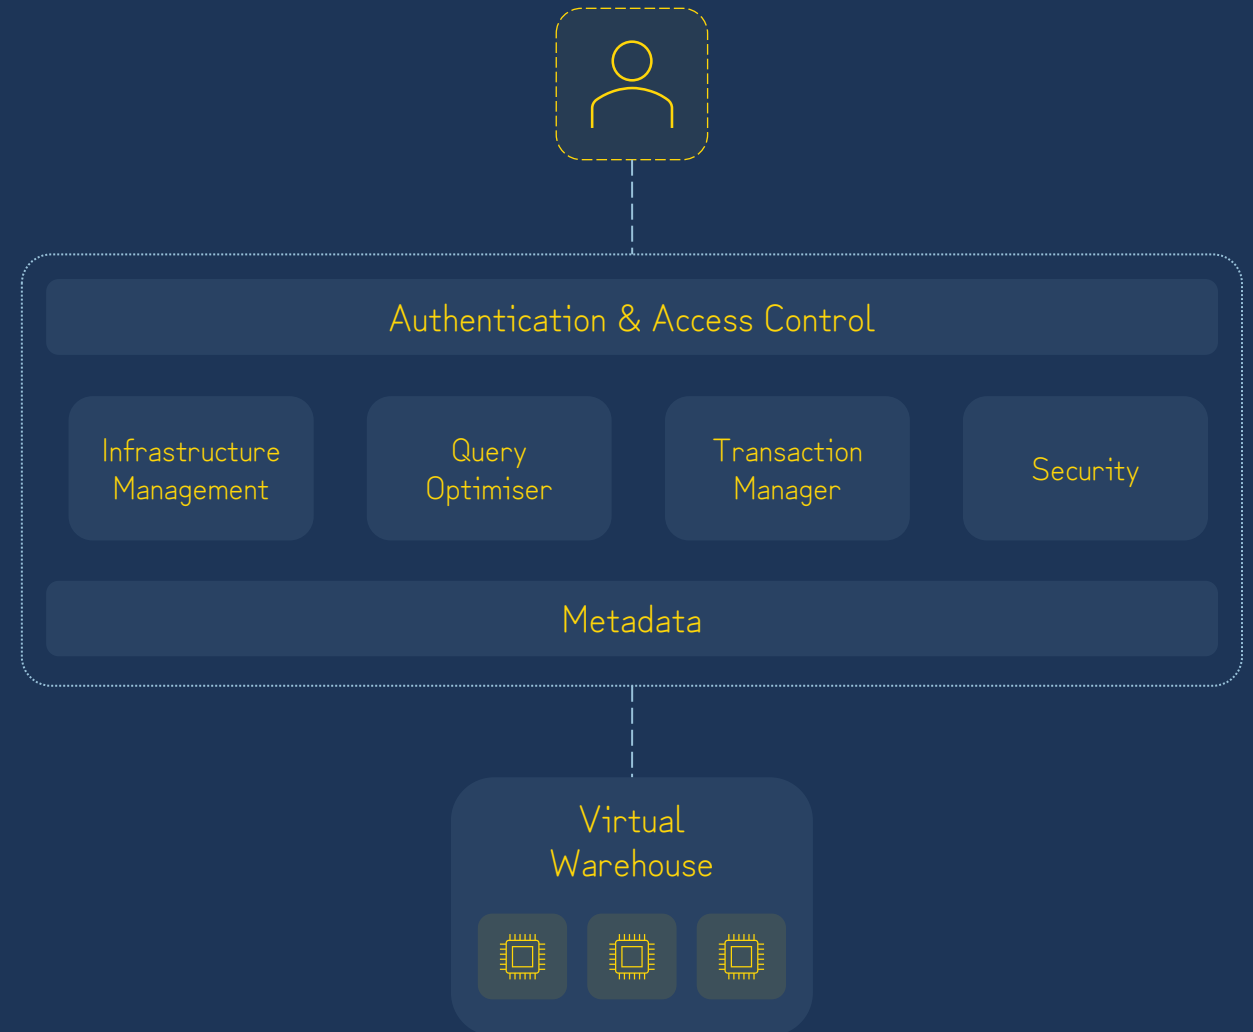
# Services Layer

The services layer is a collection of highly available and scalable services that coordinate activities such as authentication and query optimization across all Snowflake accounts.

Similar to the underlying virtual warehouse resources, the services layer also runs on cloud compute instances.

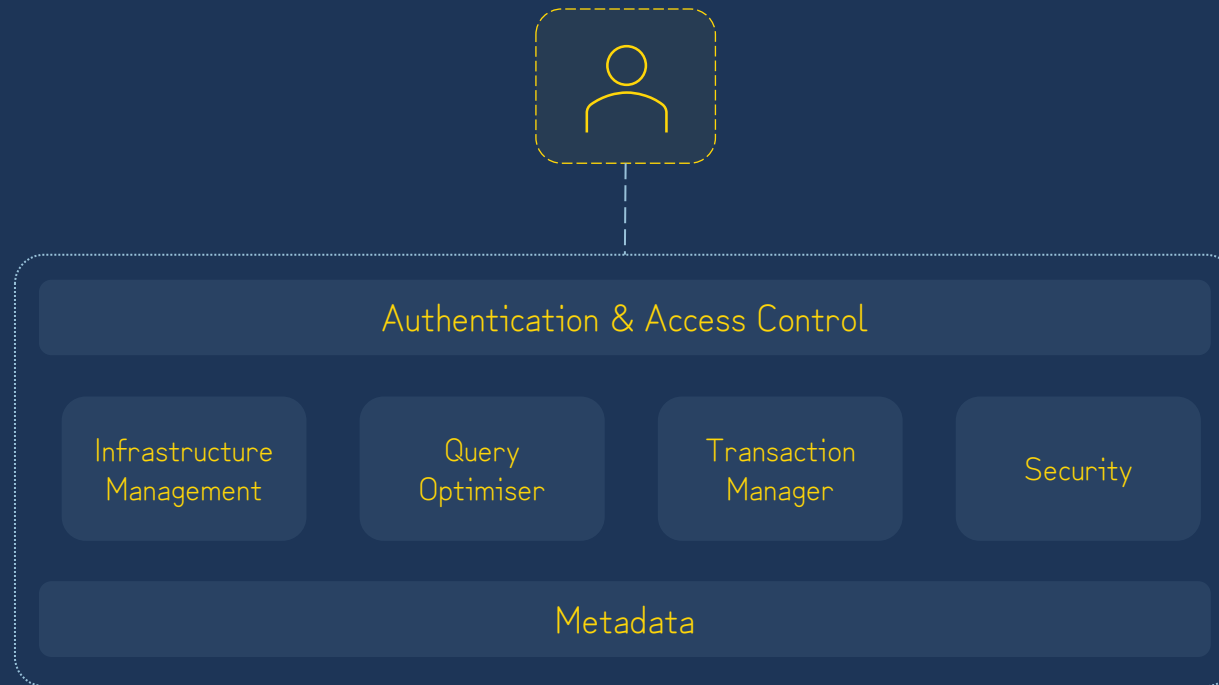
Services managed by this layer include:

- Authentication & Access Control
- Infrastructure Management
- Transaction Management
- Metadata Management
- Query parsing and optimisation
- Security



# Services Layer

The services layer is a collection of highly available and scalable services that coordinate activities such as authentication and query optimization across all Snowflake accounts.



# Editions & Key Features

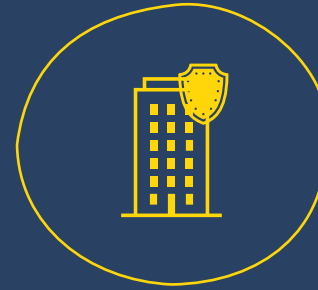
# Snowflake Editions & Key Features



Standard



Enterprise



Business Critical



Virtual Private  
Snowflake

SQL Support

Security, Governance, & Data Protection

Compute Resource Management

Interface & Tools

Releases

Data Import & Export

Data Replication & Failover

# SQL Support



Feature



Standard



Enterprise



Business Critical



VPS

⇒ Standard SQL defined in SQL:1999



⇒ Advanced DML defined in SQL:2003



⇒ Standard data types: VARCHAR, NUMBER, TIMESTAMP etc



⇒ Semi-structured data types: VARIANT, OBJECT & ARRAY.



⇒ Multi-statement transactions



**TOM BAILEY COURSES**

[tombaileycourses.com](http://tombaileycourses.com)

# SQL Support



Feature



Standard



Enterprise



Business Critical



VPS

⇒ User Defined Functions (UDFs)



⇒ Automatic Clustering



⇒ Zero-copy Cloning



⇒ Search Optimization Service



⇒ Materialized Views



# Security, Governance & Data Protection



Feature



Standard



Enterprise



Business Critical



VPS

⇒ Federated authentication and SSO



⇒ Multi-factor authentication (MFA)



⇒ Time Travel & Fail-safe



⇒ Encryption at-rest and in-transit



⇒ Network Policies



⇒ Access Control Framework



Continuous  
Data  
Protection

**TOM BAILEY COURSES**

[tombaileycourses.com](https://tombaileycourses.com)



# Security, Governance & Data Protection



Feature



Standard



Enterprise



Business Critical



VPS

⇒ Column and row access policies



⇒ Tri-secret secure



⇒ Private Connectivity



⇒ Support for compliance regulations: PCI DSS,  
HIPAA, HITRUST CSF, IRAP, FedRAMP



⇒ Dedicated metastore and pool of compute  
resources



# Compute Resource Management



Feature



Standard



Enterprise



Business Critical



VPS

⇒ Separate compute clusters (Virtual Warehouses)



⇒ Resource Monitors



⇒ Multi-cluster Virtual Warehouses



# Interface & Tools



Feature



Standard



Enterprise



Business Critical



VPS

⇒ Classic UI & Snowsight UI



⇒ SnowSQL & SnowCD



⇒ Native programming interfaces



⇒ Third-party tools ecosystem



⇒ Snowflake Partner Connect



# Data Import & Export



Feature



Standard



Enterprise



Business Critical



VPS

⇒ Bulk Loading



⇒ Bulk Unloading



⇒ Continuous Data Loading with Snowpipe



⇒ Snowflake Connector for Kafka



# Data Replication & Failover



Feature



Standard



Enterprise



Business Critical



VPS

⇒ Database replication

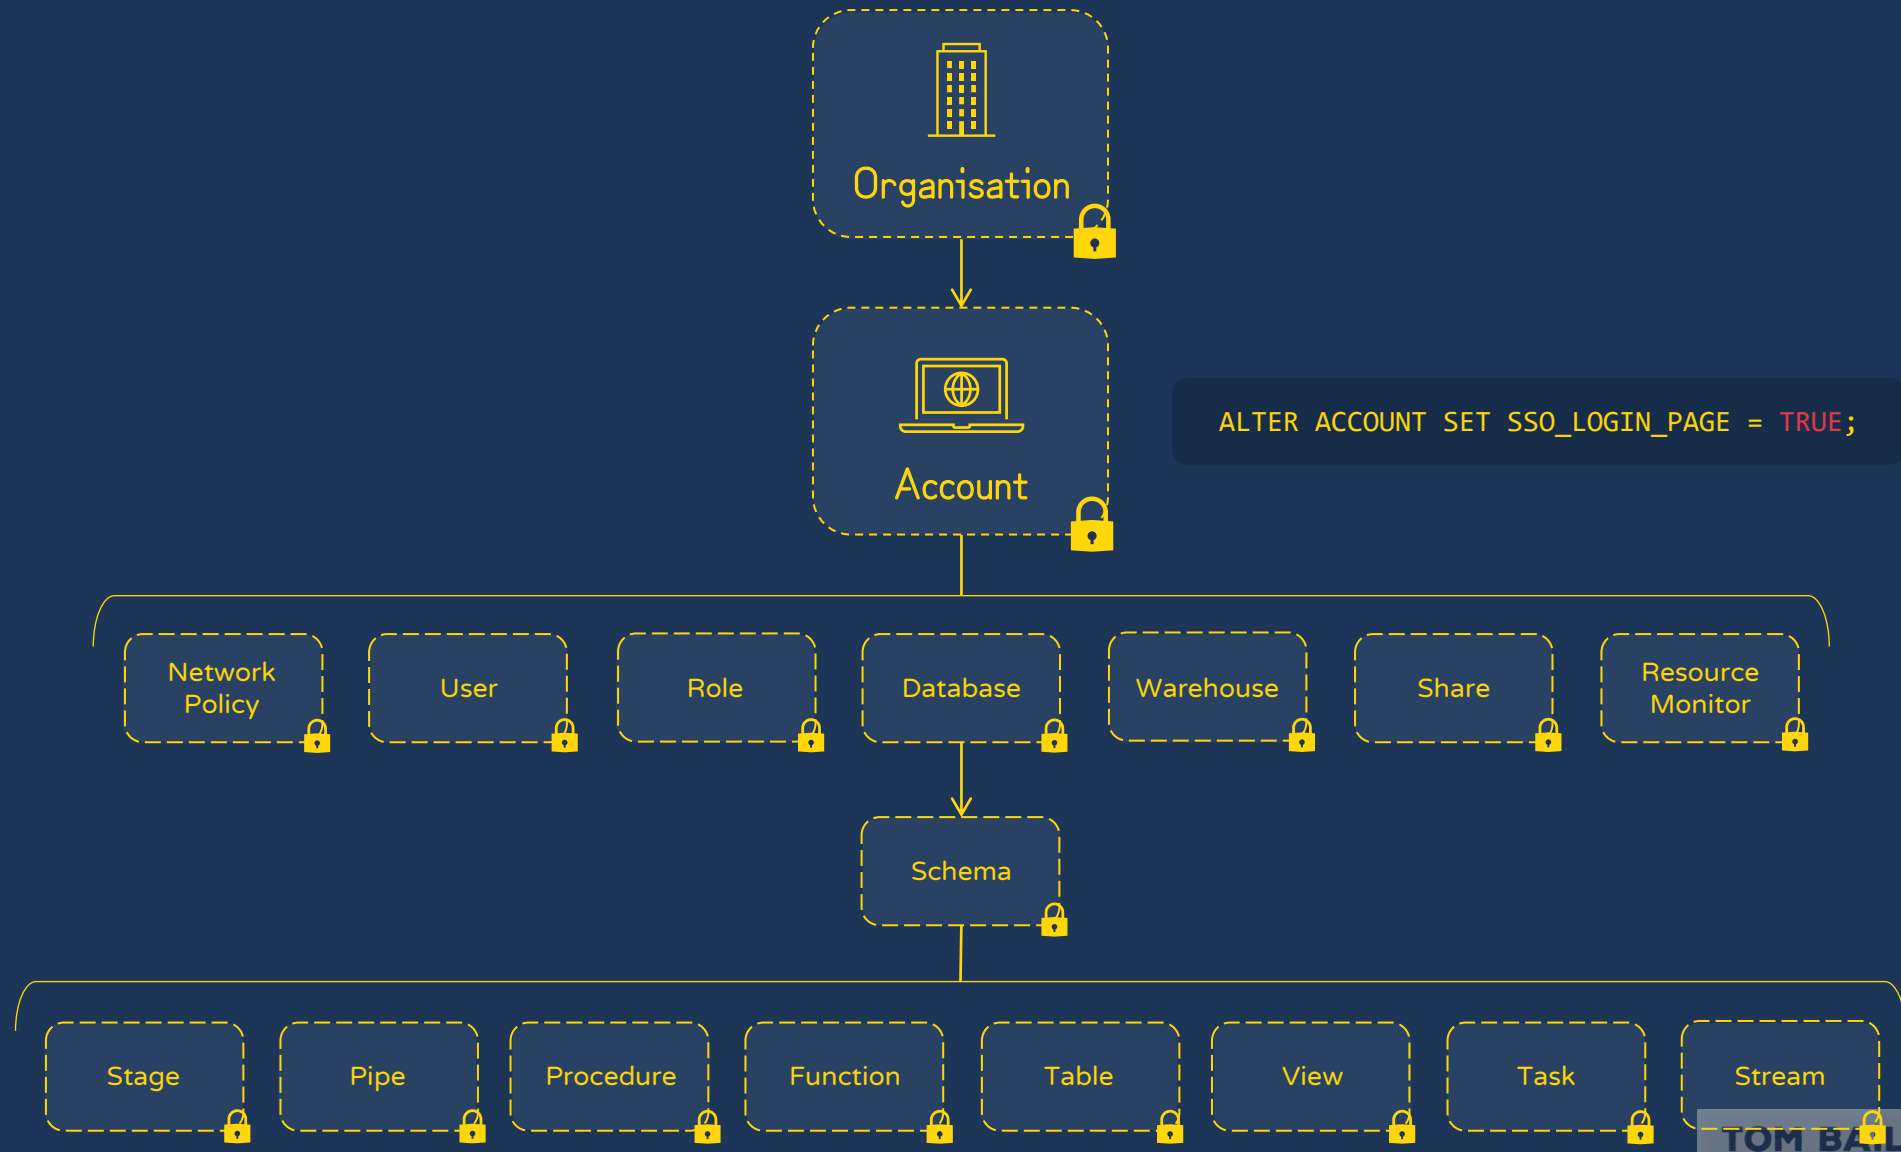


⇒ Database failover and failback



# Snowflake's Catalogue and Objects

# Snowflake Object Model



# Organisation, Account, Database & Schema.



# Organisation Overview



1

Manage one or more Snowflake accounts.

2

Setup and administer Snowflake features which make use of multiple accounts.

3

Monitoring usage across accounts.

# Organisation Setup



Contact Snowflake support



Provide organisation name and nominate an account



ORGADMIN role added to nominated account

The screenshot shows the Snowflake web interface. The top navigation bar includes icons for Databases, Shares, Data Marketplace, Warehouses, Worksheets, History, Organization (highlighted with a yellow box), Preview App, Partner Connect, Help, and a user profile for ADMIN ORGADMIN. Below the navigation bar, the 'Organization' section is active, displaying 'Accounts'. A search bar shows '45 Accounts'. Below this is a table of accounts with columns: Account Name, Edition, Snowflake Region, Date Created, and Account URL. The table lists several accounts, with 'PROD\_PRIMARY' highlighted in light blue.

Account Name	Edition	Snowflake Region	Date Created	Account URL
TEST_AZURE	Standard	AZURE_SOUTHEASTASIA	6:02:36 PM	<a href="https://sfpmdemo-test_azure.snowflakecomputing.com">https://sfpmdemo-test_azure.snowflakecomputing.com</a>
TEST_GCP	Enterprise	GCP_EUROPE_WEST2	6:01:32 PM	<a href="https://sfpmdemo-test_gcp.snowflakecomputing.com">https://sfpmdemo-test_gcp.snowflakecomputing.com</a>
TEST_AWS	Enterprise	AWS_US_EAST_1	6:00:34 PM	<a href="https://sfpmdemo-test_aws.snowflakecomputing.com">https://sfpmdemo-test_aws.snowflakecomputing.com</a>
PROD_SECONDARY	Business Critical	AWS_US_WEST_2	5:58:44 PM	<a href="https://sfpmdemo-prod_secondary.snowflakecomputing.com">https://sfpmdemo-prod_secondary.snowflakecomputing.com</a>
PROD_PRIMARY	Business Critical	AWS_US_EAST_1	5:56:54 PM	<a href="https://sfpmdemo-prod_primary.snowflakecomputing.com">https://sfpmdemo-prod_primary.snowflakecomputing.com</a>
RAMEN	Standard	AZURE_EASTUS2	1/8/2021, 4:20:53 PM	<a href="https://sfpmdemo-ramen.snowflakecomputing.com">https://sfpmdemo-ramen.snowflakecomputing.com</a>

# ORGADMIN Role

ORGADMIN

Account Management

```
CREATE ACCOUNT MYACCOUNT1  
  ADMIN_NAME = admin  
  ADMIN_PASSWORD = 'Password123'  
  FIRST_NAME = jane  
  LAST_NAME = smith  
  EMAIL = 'myemail@myorg.org'  
  EDITION = enterprise  
  REGION = aws_us_west_2;
```

```
SHOW ORGANIZATION ACCOUNTS;
```

```
SHOW REGIONS;
```

Enable cross-account features

```
SELECT  
system$global_account_set_parameter(  
  'UT677AA',  
  'ENABLE_ACCOUNT_DATABASE_REPLICATION',  
  'true');
```

Monitoring account usage

```
SELECT  
  ROUND(SUM(AVERAGE_BYTES) / POWER(1024,4),2)  
FROM ORGANIZATION_USAGE.STORAGE_DAILY_HISTORY  
WHERE USAGE_DATE = CURRENT_DATE();
```

# Account Overview

An account is the administrative name for a collection of storage, compute and cloud services deployed and managed entirely on a selected cloud platform.

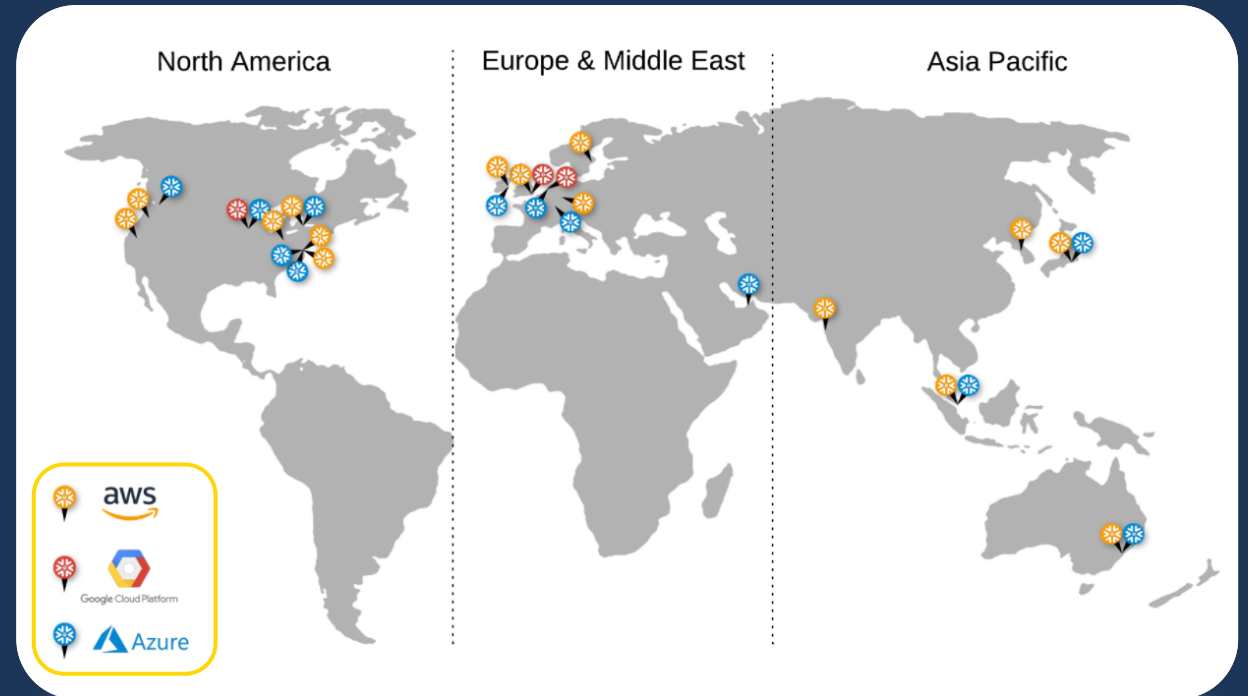
Each account is hosted on a **single cloud provider**:

- Amazon Web Services (AWS)
- Google Cloud Platform (GCP)
- Microsoft Azure (Azure)

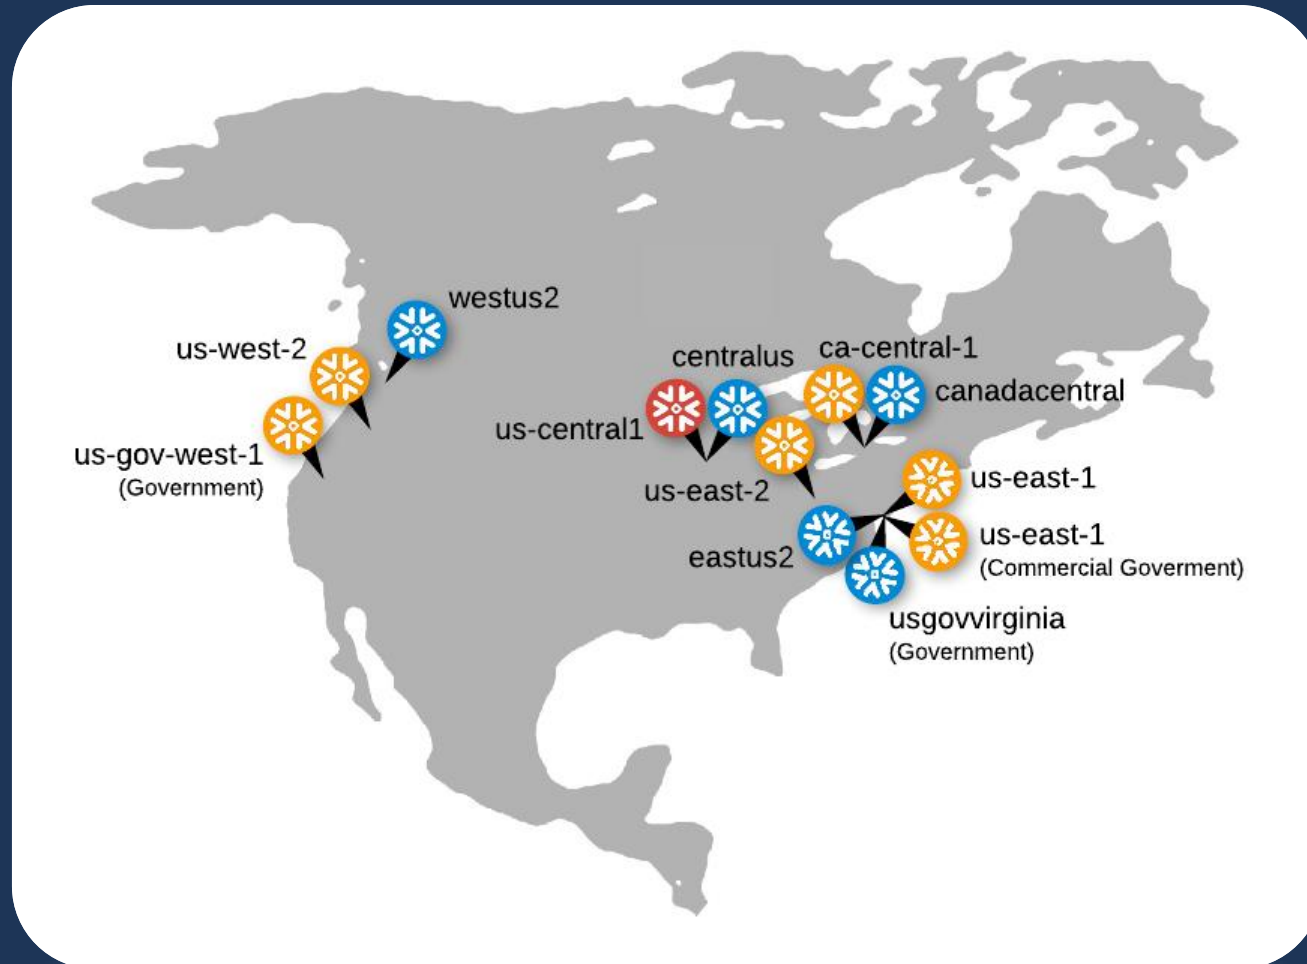
Each account is provisioned in a **single geographic region**.

Each account is created as a **single Snowflake edition**.

An account is created with the system-defined role **ACCOUNTADMIN**.



# Account Regions



aws.us-west-2

US West (Oregon)

aws.ca-central-1

Canada (Central)

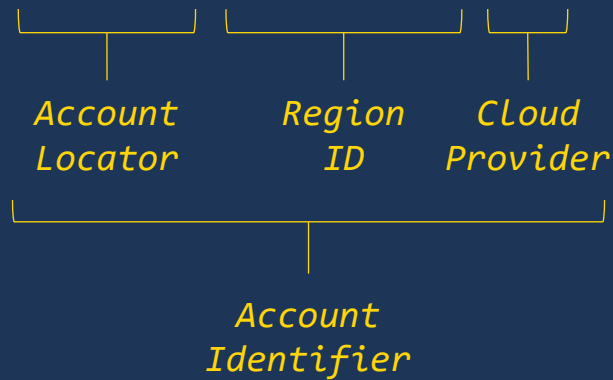
azure.westus2

West US 2 (Washington)

# Account URL

Using an Account Locator as an Identifier

xy12345.us-east-2.aws.snowflakecomputing.com

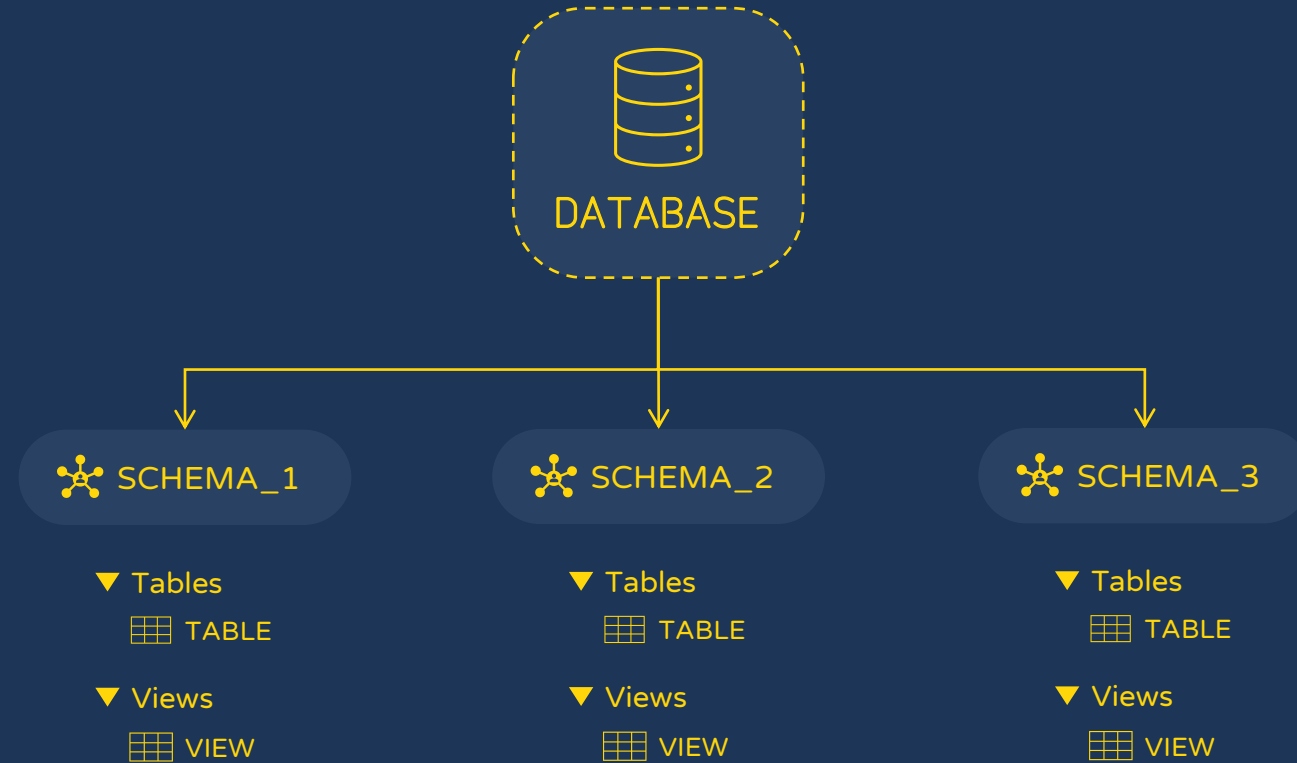


Using an Organization and Account Name as Identifier

acme-marketing-test-account.snowflakecomputing.com



# Database & Schemas



# Database & Schemas



## DATABASE

Databases must have a unique identifier in an account.

A database must start with an alphabetic character and cannot contain spaces or special characters unless enclosed in double quotes.

```
CREATE DATABASE MY_DATABASE;
```

```
CREATE DATABASE MY_DB_CLONE CLONE MYTESTDB;
```

```
CREATE DATABASE MYDB1  
  AS REPLICATION OF MYORG.ACCOUNT1.MYDB1  
  DATA_RETENTION_TIME_IN_DAYS = 10;
```

```
CREATE DATABASE SHARED_DB FROM SHARE UTT783.SHARE;
```



## SCHEMA

Schemas must have a unique identifier in a database.

A schema must start with an alphabetic character and cannot contain spaces or special characters unless enclosed in double quotes.

```
CREATE SCHEMA MY_SCHEMA;
```

```
CREATE SCHEMA MY_SCHEMA_CLONE CLONE MY_SCHEMA;
```

MY\_DATABASE.MY\_SCHEMA

*Namespace*



# Table and View Types

# Table Types



Permanent

Default table type.

Exists until explicitly dropped.



Temporary

Used for transitory data.

Persist for duration of a session.



Transient

Exists until explicitly dropped.

No fail-safe period.



External

Query data outside Snowflake.

Read-only table.

Time Travel



90 days



1 day



1 day



Fail-safe



# View Types



Standard

```
CREATE VIEW MY_VIEW AS  
SELECT COL1, COL2 FROM MY_TABLE;
```

Does not contribute to storage cost.

If source table is dropped, querying view returns error.

Used to restrict contents of a table.



Materialized

```
CREATE MATERIALIZED VIEW MY_VIEW AS  
SELECT COL1, COL2 FROM MY_TABLE;
```

Stores results of a query definition and periodically refreshes it.

Incurs cost as a serverless feature.

Used to boost performance of external tables.



Secure

```
CREATE SECURE VIEW MY_VIEW AS  
SELECT COL1, COL2 FROM MY_TABLE;
```

Both standard and materialized can be secure.

Underlying query definition only visible to authorized users.

Some query optimizations bypassed to improve security.

# UDFs and Stored Procedures

# User Defined Functions (UDFs)

User defined functions (UDFs) are schema-level objects that enable users to write their own functions in three different languages:

- SQL
- JavaScript
- Python
- Java

UDFs accept 0 or more parameters.

UDFs can return scalar or tabular results (UDTF).

UDFs can be called as part of a SQL statement.

UDFs can be overloaded.

```
CREATE FUNCTION AREA_OF_CIRCLE(radius FLOAT)
  RETURNS TABLE (area number)
AS
$$
    pi() * radius * radius
$$;
```

```
SELECT  AREA_OF_CIRCLE(col1) FROM MY_TABLE;
```

# JavaScript UDF

```
CREATE FUNCTION JS_FACTORIAL(d double)
  RETURNS DOUBLE
  LANGUAGE JAVASCRIPT
  AS
  $$
  if (D <= 0) {
    return 1
  } else {
    var result = 1;
    for (var i = 2; i <= D; i++) {
      result = result * i;
    }
    return result;
  }
  $$;
```

JavaScript is specified with the language parameter.

Enables use of high-level programming language features.

JavaScript UDFs can refer to themselves recursively.

Snowflake data types are mapped to JavaScript data types.

JavaScript

# Java UDF

```
CREATE FUNCTION DOUBLE(X INTEGER)
  RETURNS INTEGER
  LANGUAGE JAVA
  HANDLER='TestDoubleFunc.double'
  TARGET_PATH='@~/TestDoubleFunc.jar'
AS
$$
    class TestDoubleFunc {
        public static int double(int x) {
            return x * 2;
        }
    }
$$;
```

Java

Snowflake boots up a JVM to execute function written in Java.

Snowflake currently supports writing UDFs in Java versions 8.x, 9.x, 10.x, and 11.x.

Java UDFs can specify their definition as in-line code or a pre-compiled jar file.

Java UDFs cannot be designated as secure.

# External Functions

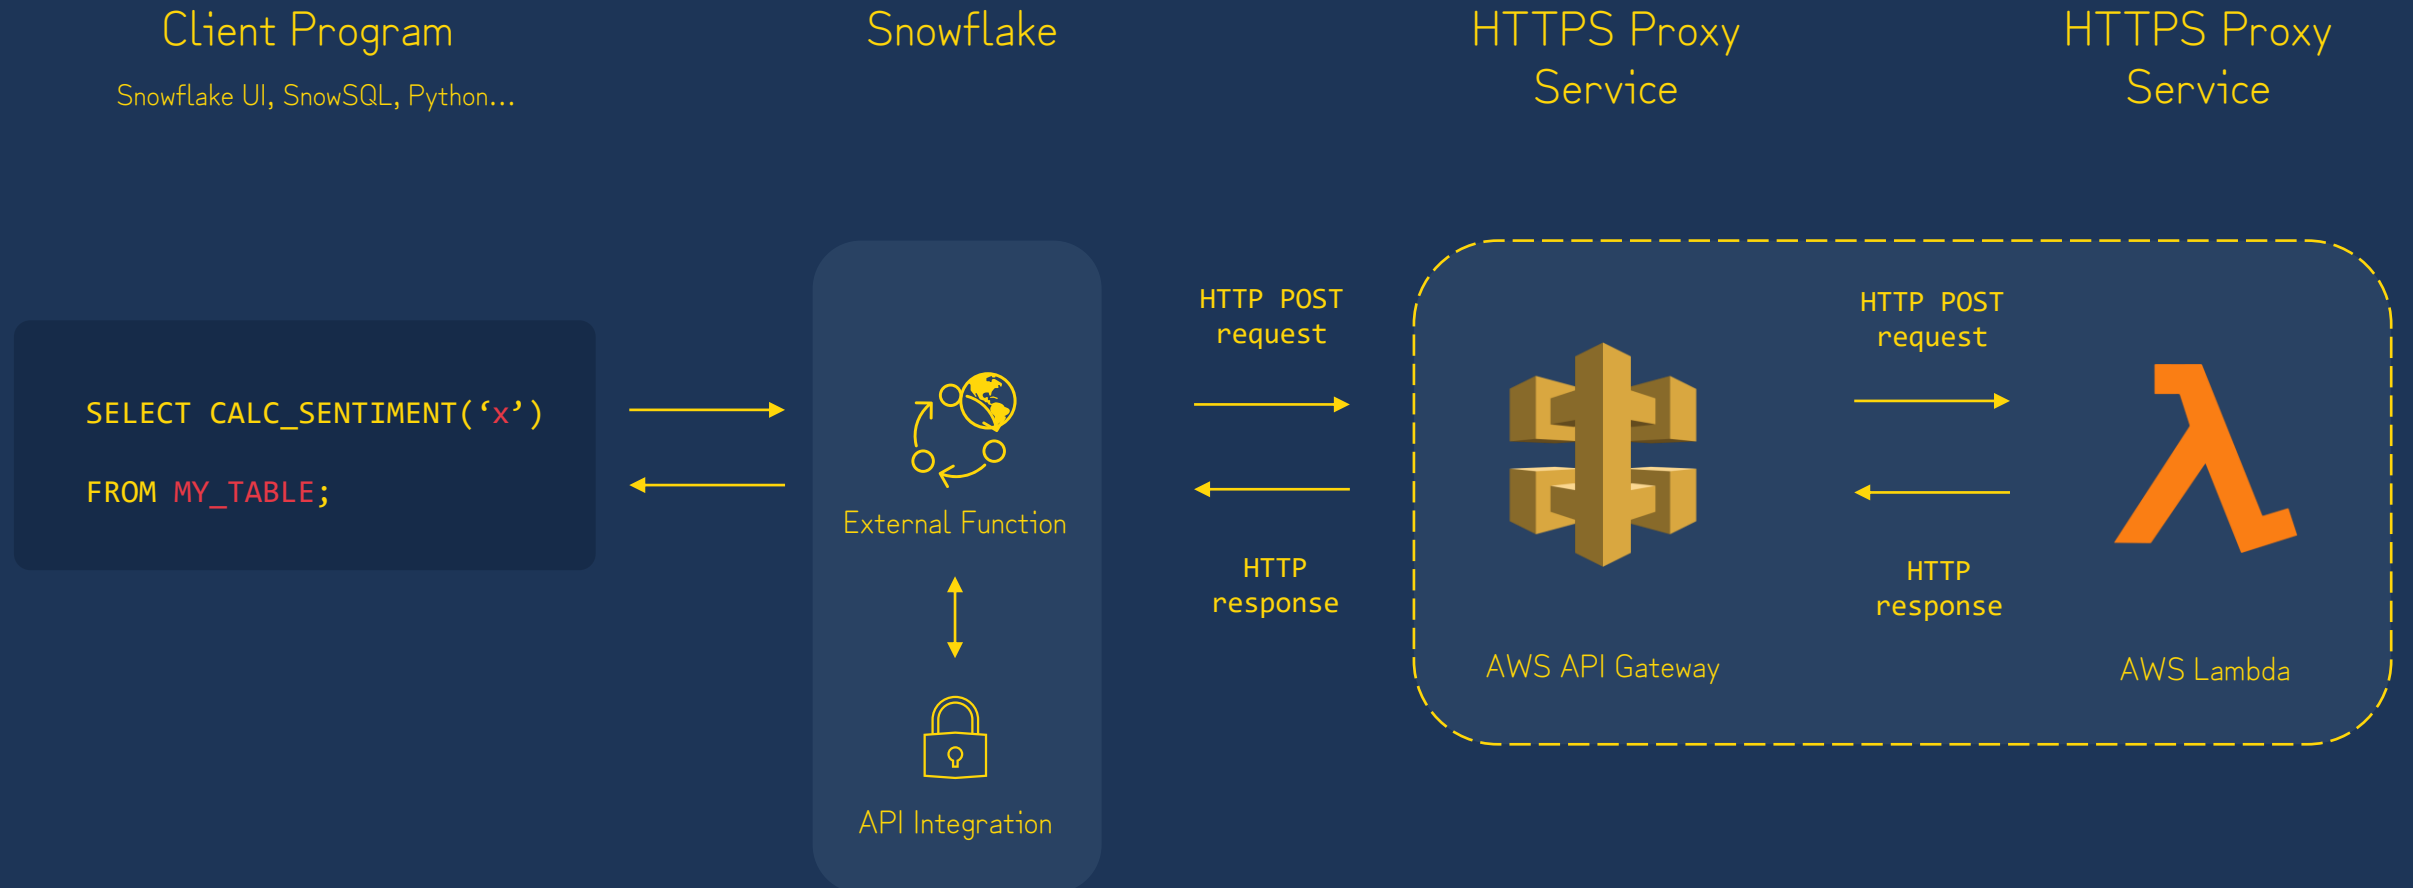
```
CREATE OR REPLACE EXTERNAL FUNCTION CALC_SENTIMENT(String_Col VARCHAR)
    RETURNS VARIANT
    API_INTEGRATION = AWS_API_INTEGRATION
    AS 'https://ttu.execute-api.eu-west-2.amazonaws.com/';
```

- ← Function Name and Parameters
- ← Return Type
- ← Integration Object
- ← URL Proxy Service

```
CREATE OR REPLACE API INTEGRATION AWS_API_INTEGRATION
    API_PROVIDER=AWS_API_GATEWAY
    API_AWS_ROLE_ARN='ARN:AWS:IAM::123456789012:ROLE/MY_CLOUD_ACCOUNT_ROLE'
    API_ALLOWED_PREFIXES=('HTTPS://XYZ.EXECUTE-API.US-WEST-2.AMAZONAWS.COM/PRODUCTION')
    ENABLED=TRUE;
```



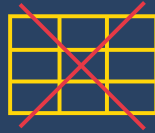
# External Function Call Lifecycle



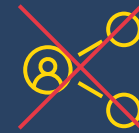
# External Function Limitations



Slower



Scalar only



Not sharable



Less secure



Egress charges

# Stored Procedures

In Relational Database Management Systems (RDBMS) **stored procedures** were **named collections of SQL statements** often containing procedural logic.



Database Admin (DBA)

```
CREATE PROCEDURE CLEAR_EMP_TABLES
DELETE FROM EMP01 WHERE EMP_DATE < DATEADD(MONTH, -1, GET_DATE())
AS
DELETE FROM EMP02 WHERE EMP_DATE < DATEADD(MONTH, -1, GET_DATE())
BEGIN
DELETE FROM EMP03 WHERE EMP_DATE < DATEADD(MONTH, -1, GET_DATE())

DELETE FROM EMP04 WHERE EMP_DATE < DATEADD(MONTH, -1, GET_DATE())
DELETE FROM EMP05 WHERE EMP_DATE < DATEADD(MONTH, -1, GET_DATE())
```

END



Data Engineer (DE)

```
EXECUTE CLEAR_EMP_TABLES;
```

 SQL examples on this slide from Microsoft SQL Server.

**TOM BAILEY COURSES**

[tombaileycourses.com](https://tombaileycourses.com)

# Snowflake Stored Procedures

①



JavaScript

②



Snowflake Scripting



SQL

③



Snowpark



Python

Java

Scala

# Stored Procedure: JavaScript

Stored procedure identifier and input parameters. —————>

RETURNS option mandatory. —————>

JAVASCRIPT, SQL, PYTHON, JAVA & SCALA. —————>

Stored procedures can execute with the owner's rights or caller's rights. —————>

Stored procedures mix JavaScript and SQL in their definition using Snowflake's JavaScript API. —————>

```
CALL EXAMPLE_STORED_PROCEDURE('EMP01');
```

```
CREATE PROCEDURE EXAMPLE_STORED_PROCEDURE(PARAM1 STRING)

    RETURNS STRING

    LANGUAGE JAVASCRIPT

    EXECUTE AS OWNER

    AS

    $$

        var param1 = PARAM1;

        var sql_command = "SELECT * FROM " + param1;

        snowflake.execute({sqlText: sql_command});

        return "Succeeded.";

    $$;
```

# Stored Procedures & UDFs

Feature	UDF	Stored Procedure
Called as part of SQL statement	✓	✗
Ability to overload	✓	✓
0 or more input parameters	✓	✓
Use of JavaScript API	✗	✓
Return of value optional	✗	✓
Values returned usable in SQL	✓	✗
Call itself recursively	✗	✓

Functions calculate something and return a value to the user.

Stored procedures perform actions rather than return values.

# Sequences

# Sequences

```
CREATE SEQUENCE DEFAULT_SEQUENCE  
  
START = 1  
  
INCREMENT = 1;
```

```
SELECT DEFAULT_SEQUENCE.NEXTVAL;
```

NEXTVAL
1

```
SELECT DEFAULT_SEQUENCE.NEXTVAL;
```

NEXTVAL
2

```
SELECT DEFAULT_SEQUENCE.NEXTVAL;
```

NEXTVAL
3



Values generated by a sequence are globally unique.



# Sequences

```
CREATE SEQUENCE INCREMENT_SEQUENCE  
  
START = 0  
  
INCREMENT = 5;
```

```
SELECT INCREMENT_SEQUENCE.NEXTVAL, INCREMENT_SEQUENCE.NEXTVAL, INCREMENT_SEQUENCE.NEXTVAL, INCREMENT_SEQUENCE.NEXTVAL;
```

NEXTVAL	NEXTVAL_1	NEXTVAL_2	NEXTVAL_3
30	40	40	50



Sequences cannot guarantee their values will be gap free.

# Sequences

① INSERT INTO TABLE.

```
CREATE SEQUENCE TRANSACTION_SEQ  
START = 1001  
INCREMENT = 1;
```

```
INSERT INTO TRANSACTION (ID)  
VALUES (TRANSACTION_SEQ.NEXTVAL)
```

```
SELECT ID FROM TRANSACTION;
```

NEXTVAL
1001

# Sequences

②

DEFAULT VALUE FOR  
A COLUMN TABLE.

```
CREATE TABLE TRANSACTIONS  
(ID INTEGER DEFAULT TRANSACTION_SEQ.NEXTVAL,  
AMOUNT DOUBLE);
```

```
INSERT INTO TRANSACTION (AMOUNT) VALUES (756.00);
```

```
SELECT ID FROM TRANSACTION;
```

ID	AMOUNT
1002	756.00

# Tasks & Streams

# Tasks & Streams



Tasks



Streams

# Tasks

A task is an object used to schedule the execution of a SQL command or a stored procedure.

## Task Workflow

① ACCOUNTNADMIN role or CREATE TASK privilege.

②

```
CREATE TASK T1
  WAREHOUSE = MYWH
  SCHEDULE = '30 MINUTE'
AS
  COPY INTO MY_TABLE
  FROM $MY_STAGE;
```

← Task Name

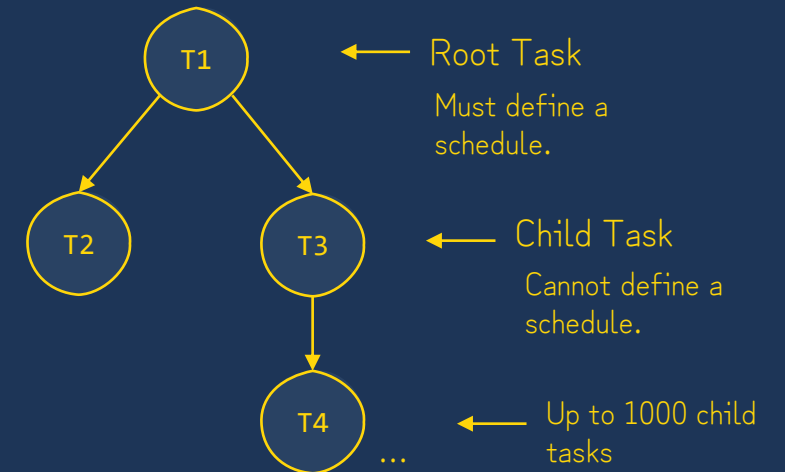
← Warehouse Definition

← Triggering Mechanism

← Query Definition

③ ALTER TASK MINUTE\_TASK RESUME; ← Start task

## Tree of Tasks



```
CREATE TASK T2
  WAREHOUSE = MYWH
  AFTER T1
AS
  COPY INTO MY_TABLE FROM $MY_STAGE;
```

← Child task triggering mechanism

# Streams

A stream is an object created to view & track DML changes to a source table – inserts, updates & deletes.

Create Stream

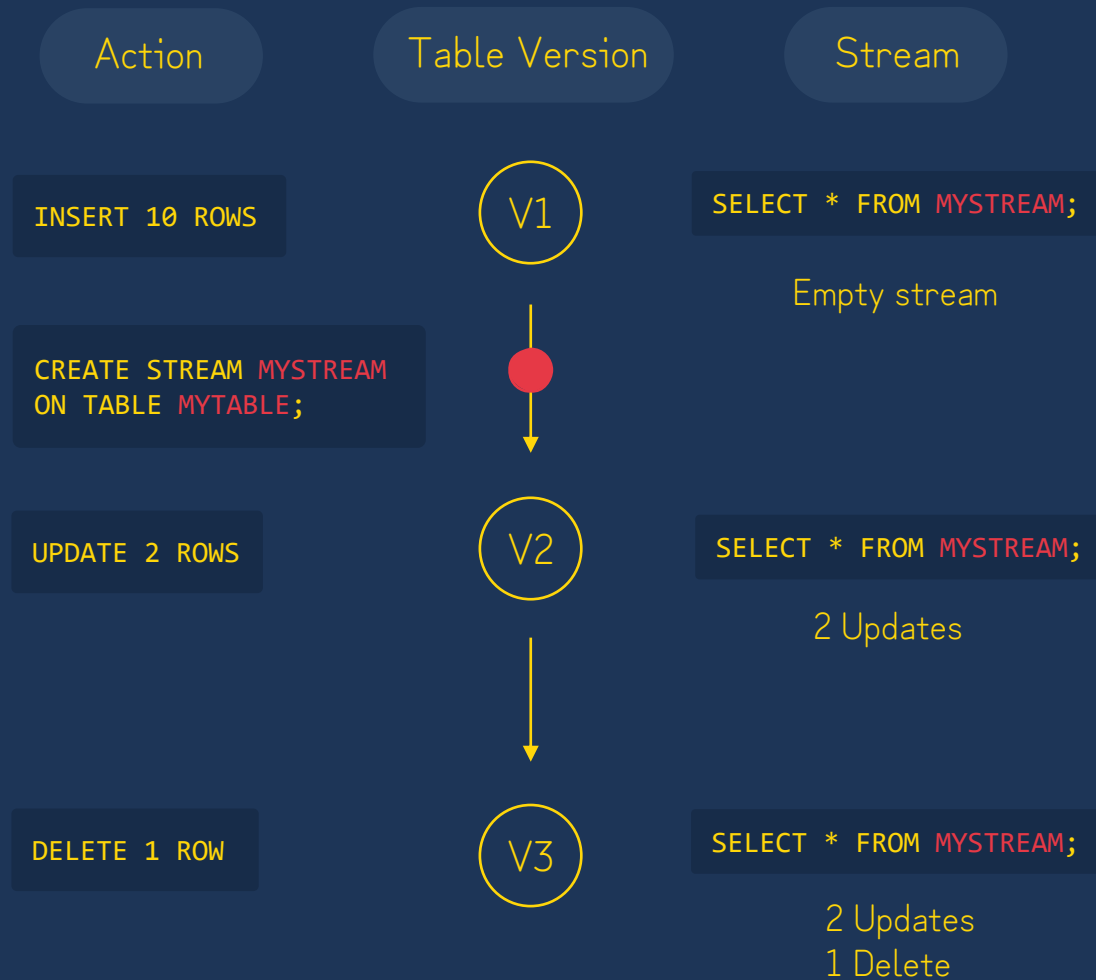
```
CREATE STREAM MY_STREAM ON TABLE MY_TABLE;
```

Query Stream

```
SELECT * FROM MY_STREAM;
```

EMP_ID	EMP_NAME	EMP_DOB	EMP_POSITION	METADATA\$ACTION	METADATA\$ISUPDATE	METADATA\$ROW_ID
AC1109	Ramesh Aravind	10/09/1964	Actor	INSERT	FALSE	cc576bf4fee43b88c4fd03

# Streams



## Progress stream offset

```
INSERT INTO MYTABLE2 SELECT * FROM MYSTREAM;
```

## Tasks & Streams

```
CREATE TASK MYTASK1
  WAREHOUSE = MYWH
  SCHEDULE = '5 MINUTE'
  WHEN
    SYSTEM$STREAM_HAS_DATA('MYSTREAM')
  AS
    INSERT INTO MYTABLE1(ID,NAME) SELECT ID, NAME
    FROM MYSTREAM WHERE METADATA$ACTION = 'INSERT';
```



# Billing

# Billing Overview



On-demand

Pay for usage as you go



Capacity

Pay for usage upfront


# Billing Overview



Virtual Warehouse  
Services



Cloud Services



Serverless Services



Storage




Data Transfer

# Billing Overview



Credits



Dollar Value



# Compute Billing Overview



**Credits** | Snowflake billing unit of measure for compute resource consumption.



## Virtual Warehouse Services

- Credit calculated based on size of virtual warehouse.
- Credit calculated on per second basis while a virtual warehouse is in 'started' state.
- Credit calculated with a minimum of 60 seconds.



## Cloud Services

- Credits calculated at a rate of 4.4 Credits per compute hour.
- Only cloud services that exceeds 10% of the daily usage of the compute resources are billed.
- This is called the Cloud Services Adjustment.



## Serverless Services

- Each serverless feature has it's own credit rate per compute-hour.
- Serverless features are composed of both compute services and cloud services.
- Cloud Services Adjustment does not apply to cloud services usage when used by serverless features.

# Data Storage & Transfer Billing Overview

\$ Dollar value | Storage and Data Transfer are billed in currency.



Data Storage



Data Transfer

- Data storage is calculated monthly based on the average number of on-disk bytes per day in the following locations:
  - Database Tables.
  - Internal Stages.
- Costs calculated based on a flat dollar value rate per terabyte (TB) based on:
  - Capacity or On-demand.
  - Cloud provider.
  - Region.
- Data transfer charges apply when moving data from one region to another or from one cloud platform to another.
- Unloading data from Snowflake using **COPY INTO <location>** command.
- Replicating data to a Snowflake account in a different region or cloud platform.
- External functions transferring data out of and into Snowflake.

# SnowCD

①

```
SELECT SYSTEM$WHITELIST();
```

Returns hostnames and port numbers.



```
1 [{"type":"SNOWFLAKE_DEPLOYMENT","host":"ht09440.eu-west-2.aws.snowflakecomputing.com","port":443},
2 {"type":"SNOWFLAKE_DEPLOYMENT_REGIONLESS","host":"jilikhy-ke89319.snowflakecomputing.com","port":443},
3 {"type":"STAGE","host":"sfc-uk-ds1-3-customer-stage.s3.eu-west-2.amazonaws.com","port":443},
4 {"type":"STAGE","host":"sfc-uk-ds1-3-customer-stage.s3.eu-west-2.amazonaws.com","port":443},
5 {"type":"STAGE","host":"sfc-uk-ds1-3-customer-stage.s3.amazonaws.com","port":443},
6 {"type":"SNOWSQL_REPO","host":"sfc-repo.snowflakecomputing.com","port":443},
7 {"type":"OUT_OF_BAND_TELEMETRY","host":"client-telemetry.snowflakecomputing.com","port":443},
8 {"type":"OCSP_CACHE","host":"ocsp.snowflakecomputing.com","port":80},
9 {"type":"DUO_SECURITY","host":"api-edddc45f.duosecurity.com","port":443},
10 {"type":"OCSP_RESPONDER","host":"ocsp.rootg2.amazontrust.com","port":80},
11 {"type":"OCSP_RESPONDER","host":"o.ss2.us","port":80},
12 {"type":"OCSP_RESPONDER","host":"ocsp.sca1b.amazontrust.com","port":80},
13 {"type":"OCSP_RESPONDER","host":"ocsp.rootca1.amazontrust.com","port":80}]
```

```
13 [{"type":"OCSP_RESPONDER","host":"ocsp.rootg2.amazontrust.com","port":80},
14 {"type":"OCSP_RESPONDER","host":"ocsp.sca1b.amazontrust.com","port":80},
15 {"type":"OCSP_RESPONDER","host":"o.ss2.us","port":80}]
```

# SnowCD

whitelist.json

```
1 [{"type": "SNOWFLAKE_DEPLOYMENT", "host": "ht09440.eu-west-2.aws.snowflakecomputing.com", "port": 443},
2 {"type": "SNOWFLAKE_DEPLOYMENT_REGIONLESS", "host": "jilikh-y-ke89319.snowflakecomputing.com", "port": 443},
3 {"type": "STAGE", "host": "sfc-uk-ds1-3-customer-stage.s3.eu-west-2.amazonaws.com", "port": 443},
4 {"type": "STAGE", "host": "sfc-uk-ds1-3-customer-stage.s3.eu-west-2.amazonaws.com", "port": 443},
5 {"type": "STAGE", "host": "sfc-uk-ds1-3-customer-stage.s3.amazonaws.com", "port": 443},
6 {"type": "SNOWSQL_REPO", "host": "sfc-repo.snowflakecomputing.com", "port": 443},
7 {"type": "OUT_OF_BAND_TELEMETRY", "host": "client-telemetry.snowflakecomputing.com", "port": 443},
8 {"type": "OCSP_CACHE", "host": "ocsp.snowflakecomputing.com", "port": 80},
9 {"type": "DUO_SECURITY", "host": "api-edddc45f.duosecurity.com", "port": 443},
10 {"type": "OCSP_RESPONDER", "host": "ocsp.rootg2.amazontrust.com", "port": 80},
11 {"type": "OCSP_RESPONDER", "host": "o.ss2.us", "port": 80},
12 {"type": "OCSP_RESPONDER", "host": "ocsp.scalb.amazontrust.com", "port": 80},
13 {"type": "OCSP_RESPONDER", "host": "ocsp.rootcal.amazontrust.com", "port": 80}]
14 [{"type": "OCSP_RESPONDER", "host": "ocsp.snowflakecomputing.com", "port": 80}]
15 [{"type": "OCSP_RESPONDER", "host": "ocsp.snowflakecomputing.com", "port": 80}]
16 [{"type": "OCSP_RESPONDER", "host": "ocsp.snowflakecomputing.com", "port": 80}]
```



snowcd ~\whitelist.json



Performing 30 checks on 12 hosts  
All checks passed.



Check for 5 hosts failed, display as follow:

```
=====
Host: ocsp.snowflakecomputing.com
Port: 80
Type: OCSP_CACHE
Failed Check: HTTP checker
Error: Invalid http code received: 400 Bad Request
Suggestion: Check the connection to your http host or transparent Proxy
```



# Connectivity: Connectors, Drivers and Partnered Tools

# Connectors and Drivers



Python



Go



PHP



.NET



NodeJS



Spark



Kafka



JDBC



ODBC

ODBC

Connect third-party tools

TOM BAILEY COURSES

[tombaileycourses.com](https://tombaileycourses.com)

# Python Connector Example

```
pip install snowflake-connector-python==2.6.2
```

```
#!/usr/bin/env python
import snowflake.connector

# Gets the version
ctx = snowflake.connector.connect(
    user='<user_name>',
    password='<password>',
    account='<account_identifier>'
)
cs = ctx.cursor()
try:
    cs.execute("SELECT current_version()")
    one_row = cs.fetchone()
    print(one_row[0])
finally:
    cs.close()
ctx.close()
```

# Snowflake Partner Tools



Business Intelligence



Data Integration



Security & Governance



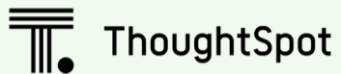
SQL Development &  
Management



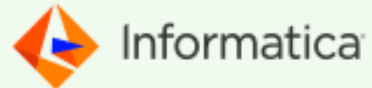
Machine Learning &  
Data Science

# Snowflake Partner Tools

## Business Intelligence



## Data Integration



## Security & Governance



 Snowflake Partner Connect is a feature to expedite connectivity with partnered tools.

TOM BAILEY COURSES

[tombaileycourses.com](https://tombaileycourses.com)

# Snowflake Partner Tools

## Machine Learning & Data Science



**DataRobot**



## SQL Development & Management



# Snowflake Scripting

# Snowflake Scripting



**Snowflake Scripting** is an extension to Snowflake SQL that adds support for **procedural logic**.

It's used to write **stored procedures** and **procedural code** outside of a stored procedure.

```
DECLARE
```

```
    (variable declarations, cursor declarations, etc.)
```

```
BEGIN
```

```
    (Snowflake Scripting and SQL statements)
```

```
EXCEPTION
```

```
    (statements for handling exceptions)
```

```
END;
```



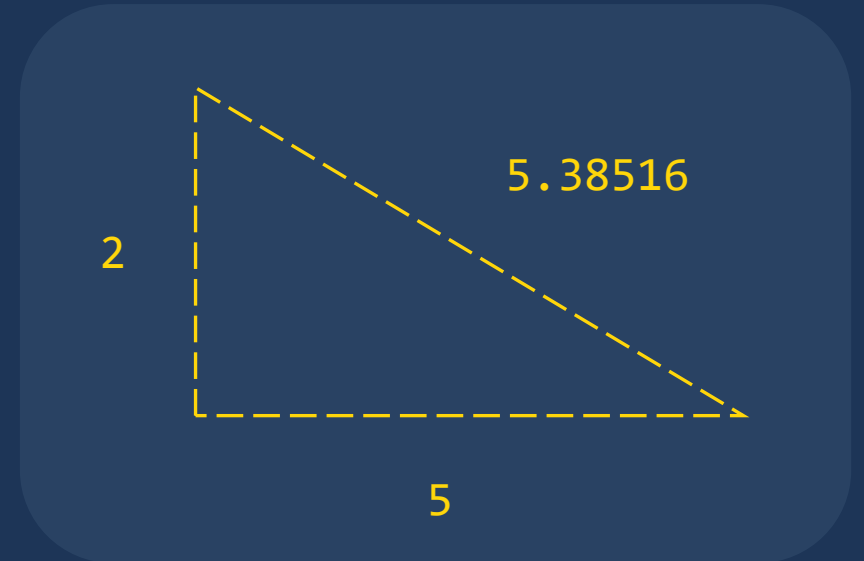
# Snowflake Scripting

```
declare
    leg_a number(38, 2);
    hypotenuse number(38,5);
begin
    leg_a := 2;
    let leg_b := 5;

    hypotenuse := sqrt(square(leg_a) + square(leg_b));
    return hypotenuse;
end;
```

anonymous block

5.38516



① Variables can only be used within the scope of the block.

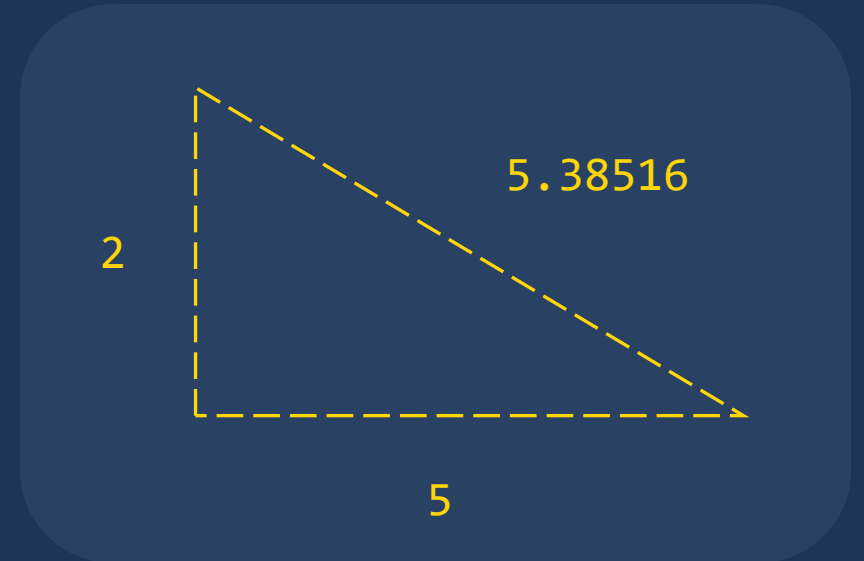
① Variables can also be declared and assigned in the BEGIN section using the LET keyword.

TOM BAILEY COURSES

tombaileycourses.com

# Snowflake Scripting

```
CREATE PROCEDURE pythagoras()  
RETURNS float  
LANGUAGE sql  
AS  
declare  
    leg_a number(38, 2);  
    hypotenuse number(38,5);  
begin  
    leg_a := 2;  
    let leg_b := 5;  
  
    hypotenuse := sqrt(square(leg_a) + square(leg_b));  
    return hypotenuse;  
end;
```



SnowSQL and the Classic Console do not correctly parse Snowflake Scripting blocks, they need to be wrapped in string constant delimiters like dollar signs.

# Branching Constructs

```
begin
  let count := 4;
  if (count % 2 = 0) then
    return 'even value';
  else
    return 'odd value';
  end if;
end;
```



DECLARE or EXCEPTION sections of a block are optional.

anonymous block
even value

# Looping Constructs

```
declare
    total integer default 0;
    max_num integer default 10;
begin
    for i in 1 to max_num do
        total := i + total;
    end for;
    return total;
end;
```

anonymous block

55

# Cursor

```
declare
    total_amount float;
    c1 cursor for select amount from transactions;
begin
    total_amount := 0.0;
    for record in c1 do
        total_amount := total_amount + record.amount;
    end for;
    return total_amount;
end;
```

anonymous block
-----------------

136.78
--------

# RESULTSET

## TABLE()

```
declare
    res resultset;
begin
    res := (select amount from transactions);
    return table(res);
end;
```

amount
101.01
24.78
10.99

# RESULTSET

## Cursor

```
declare
    total_amount float;
    res resultset default (select amount from transactions);
    c1 cursor for res;
begin
    total_amount := 0.0;
    for record in c1 do
        total_amount := total_amount + record.amount;
    end for;
    return total_amount;
end;
```

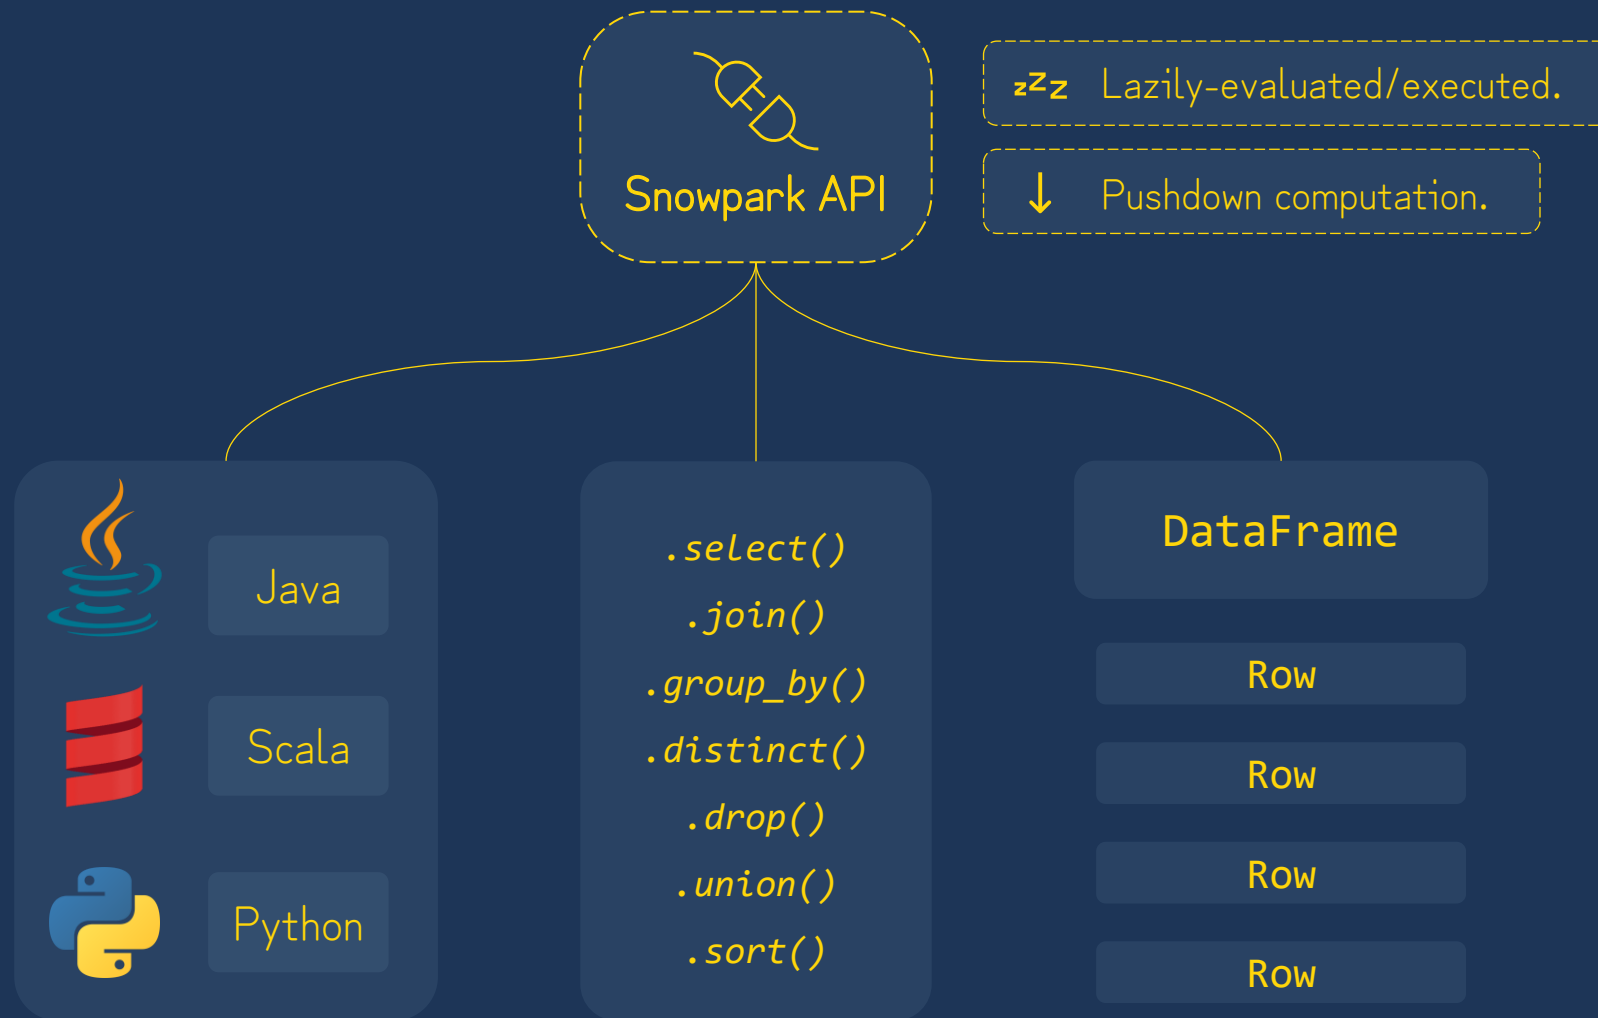
anonymous block

136.78

# Snowpark



# Snowpark



# Snowpark API: Python

```
1 import os
2 from snowflake.snowpark import Session
3 from snowflake.snowpark.functions import col
```

```
5 connection_parameters = {
6     "account": os.environ["snowflake_account"],
7     "user": os.environ["snowflake_user"],
8     "password": os.environ["snowflake_password"],
9     "role": os.environ["snowflake_user_role"],
10    "warehouse": os.environ["snowflake_warehouse"],
11    "database": os.environ["snowflake_database"],
12    "schema": os.environ["snowflake_schema"]
13 }
```

# Snowpark API: Python

```
14 session = Session.builder.configs(connection_parameters).create()

15 transactions_df = session.table("transactions")

16 print(transactions_df.collect())
```

Console output:

```
[Row(ACCOUNT_ID=8764442, AMOUNT=12.99),
Row(ACCOUNT_ID=8764442, AMOUNT=50.0),
Row(ACCOUNT_ID=8764442, AMOUNT=1100.0),
Row(ACCOUNT_ID=8764443, AMOUNT=110.0),
Row(ACCOUNT_ID=8764443, AMOUNT=2766.0),
Row(ACCOUNT_ID=8764443, AMOUNT=1010.0),
Row(ACCOUNT_ID=8764443, AMOUNT=3022.23),
Row(ACCOUNT_ID=8764444, AMOUNT=6986.0),
Row(ACCOUNT_ID=8764444, AMOUNT=1500.0)]
```

# Snowpark API: Python

```
18 transactions_df_filtered = transactions_df.filter(col("amount") >= 1000.00)

19 transaction_counts_df = transactions_df_filtered.group_by("account_id").count()

20 flagged_transactions_df = transaction_counts_df.filter(col("count") >= 2).rename(col("count"), "flagged_count")

21 flagged_transactions_df.write.save_as_table("flagged_transactions", mode="append")

22 print(flagged_transactions_df.show())

23 session.close()
```

Console output:

"ACCOUNT_ID"	"FLAGGED_COUNT"
8764443	3
8764444	2